

Load Required Packages

```
In [2]: import os,sys
import pandas as pd
import scipy.io
import numpy as np
from scipy.spatial.distance import pdist
import time
import math
import keras
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from sklearn import ensemble
from sklearn.metrics import accuracy_score, make_scorer, classification_report
from statistics import mean
from sklearn.ensemble import GradientBoostingClassifier
import PIL
from PIL import Image
from scipy.io import loadmat
from sklearn.model_selection import train_test_split, cross_validate, GridSearchCV
from keras.layers import Dense, Activation, Flatten, Input, Dropout
from keras.layers import BatchNormalization
from keras.models import Model
from keras import initializers
from keras.optimizers import Adam
from keras.utils import to_categorical
```

Using TensorFlow backend.

Part I Baseline Model: GBM

1. Provide directories for training/testing images.

```
In [15]: root = sys.path[0]
train_dir = os.path.join(root, '../data/train_set')
train_image_dir = os.path.join(train_dir, 'images')
train_pt_dir = os.path.join(train_dir, 'points')
train_label_path = os.path.join(train_dir, "label.csv")
```

2. Train/Test Split Feature Extraction

```
In [16]: info = pd.read_csv(train_label_path)
# read mat file and store coordinates in mat
m = []
for idx in info['Index']:
    file = "%04d.mat"%(idx)
    m.append( scipy.io.loadmat( os.path.join( train_pt_dir, file ) ))

mat = [x[[i for i in x.keys() if not i in ['__header__', '__version__', '__globals__']][0]] for x in m]
c = np.array([pdist(x) for x in mat[0:]])
```

```
In [17]: train_idx, test_idx = train_test_split(info['Index'], test_size=0.2, random_state=123)
```

```
In [18]: start_time_test=time.time()
train_features=np.array([pdist(mat[i-1]) for i in train_idx ])
print("baseline train features extracting takes %s seconds" % round((time.time() - start_time_test),3))
start_time_train=time.time()
test_features=np.array([pdist(mat[i-1]) for i in test_idx ])
print("base line test features extracting takes %s seconds" % round((time.time() - start_time_train),3))
train_labels=info.emotion_idx[train_idx-1]
test_labels=info.emotion_idx[test_idx-1]
print(train_features.shape,train_labels.shape)
```

```
baseline train features extracting takes 0.121 seconds
base line test features extracting takes 0.027 seconds
(2000, 3003) (2000,)
```

3. Train a Baseline GBM model with training features and responses

In [63]: *#baseline GBM*

```

baseline = GradientBoostingClassifier(learning_rate=0.1, n_estimators=100,max_
depth=3, min_samples_split=2, min_samples_leaf=1, subsample=1,max_features='sq
rt', random_state=10)
start_time=time.time()
baseline.fit(train_features, train_labels)
print("training model takes %s seconds" % round((time.time() - start_time),3
))
predictors=list(train_features)

print('Accuracy of the GBM on test set: {:.3f}'.format(baseline.score(test_fea
tures,test_labels)))
start_time1 = time.time()
pred=baseline.predict(test_features)
print("testing model takes %s seconds" % round((time.time() - start_time1),3))
print(classification_report(test_labels, pred))

```

training model takes 66.134 seconds

Accuracy of the GBM on test set: 0.440

testing model takes 0.044 seconds

	precision	recall	f1-score	support			
1	0.50	0.56	0.53	18			
2	0.65	0.68	0.67	19			
3	0.40	0.56	0.47	25			
4	0.50	0.52	0.51	21			
5	0.58	0.61	0.59	18			
6	0.63	0.46	0.53	26			
7	0.40	0.50	0.44	20			
8	0.73	0.69	0.71	16			
9	0.78	0.56	0.65	25			
10	0.45	0.45	0.45	20			
11	0.41	0.50	0.45	24			
12	0.44	0.25	0.32	32			
13	0.12	0.17	0.14	24			
14	0.54	0.57	0.55	23			
15	0.59	0.45	0.51	22			
16	0.73	0.73	0.73	22			
17	0.37	0.42	0.39	26			
18	0.31	0.23	0.26	22			
19	0.23	0.22	0.22	23			
20	0.17	0.30	0.21	20			
21	0.46	0.32	0.38	34			
22	0.25	0.20	0.22	20			
accuracy				0.44	500		
macro avg				0.47	0.45	0.45	500
weighted avg				0.46	0.44	0.44	500

4. Parameter tuning

4.1 parameter tuning Learning_rate,n_estimator

```
In [9]: #parameter tuning Learning_rate,n_estimator
#p_test3 = {'learning_rate':[0.1,0.05], 'n_estimators':[50,100,250,500,750,1000,1250,1500,1750]}

#tuning = GridSearchCV(estimator =GradientBoostingClassifier(max_depth=4, min_samples_split=2, min_samples_leaf=1, subsample=1,max_features='sqrt', random_state=10),
#
#                      param_grid = p_test3, scoring='accuracy',n_jobs=4,iid=False, cv=5)
#tuning.fit(train_features,train_labels)
#tuning.cv_results_, tuning.best_params_, tuning.best_score_
```

4.2 Tuning Max_depth, Min_samples_split

```
In [10]: #param_test2 = {'max_depth':range(1,16,2), 'min_samples_split':range(2,102,20)}
#tuning2 = GridSearchCV(estimator = GradientBoostingClassifier(learning_rate=0.05, n_estimators=500, max_features='sqrt', subsample=1, random_state=10),
#param_grid = param_test2, scoring='accuracy',n_jobs=4,iid=False, cv=5)
#tuning2.fit(train_features,train_labels)
#tuning2.cv_results_, tuning2.best_params_, tuning2.best_score_
```

4.3 Parameter Tuning:min_samples_split,min_samples_leaf

```
In [11]: #param_test3 = {'min_samples_split':range(2,102,20), 'min_samples_leaf':range(30,71,10)}
#tuning3 = GridSearchCV(estimator = GradientBoostingClassifier(learning_rate=0.05, n_estimators=500,max_depth=5,min_samples_split=62,max_features='sqrt', subsample=1, random_state=10),
#param_grid = param_test3,scoring='accuracy',n_jobs=4,iid=False, cv=5)
#tuning3.fit(train_features, train_labels)
#tuning3.cv_results_, tuning3.best_params_, tuning3.best_score_
```

4.4 Parameter Tuning Max_features

```
In [12]: #param_test4 = {'max_features':range(7,20,2)}
#tuning4 = GridSearchCV(estimator = GradientBoostingClassifier(learning_rate=0.05, n_estimators=500,max_depth=5, min_samples_split=62, min_samples_leaf=30, subsample=0.8, random_state=10),
#param_grid = param_test4, scoring='accuracy',n_jobs=4,iid=False, cv=5)
#tuning4.fit(train_features,train_labels)
#tuning4.cv_results_, tuning4.best_params_, tuning4.best_score_
```

4.5 Parameter Tuning Subsample

```
In [13]: #param_test5 = {'subsample':[0.6,0.7,0.75,0.8,0.85,0.9,1.0]}
#tuning5 = GridSearchCV(estimator = GradientBoostingClassifier(learning_rate=
0.05, n_estimators=500,max_depth=5,min_samples_split=62, min_samples_leaf=30,r
andom_state=10,max_features='sqrt'),
#param_grid = param_test5, scoring='accuracy',n_jobs=4,iid=False, cv=5)
#tuning5.fit(train_features,train_labels)
#tuning5.cv_results_, tuning5.best_params_, tuning5.best_score_
```

**5. Final Parameter set at: learning_rate=0.05,
n_estimators=500,max_depth=5,min_samples_split=62,
min_samples_leaf=30,random_state=10,max_features='sqrt',subsample=1.0**

```
In [19]: baseline_tune=GradientBoostingClassifier(learning_rate=0.05, n_estimators=500,
max_depth=5,min_samples_split=62, min_samples_leaf=30,random_state=10,max_features='sqrt',subsample=1.0)
start_time=time.time()
baseline_tune.fit(train_features, train_labels)
print("training model takes %s seconds" % round((time.time() - start_time),3))
predictors=list(train_features)

print('Accuracy of the GBM on test set: {:.3f}'.format(baseline_tune.score(test_features,test_labels)))
start_time1 = time.time()
pred=baseline_tune.predict(test_features)
print("testing model takes %s seconds" % round((time.time() - start_time1),3))
print(classification_report(test_labels, pred))
```

training model takes 282.097 seconds

Accuracy of the GBM on test set: 0.486

testing model takes 0.358 seconds

	precision	recall	f1-score	support
1	0.57	0.72	0.63	18
2	0.70	0.84	0.76	19
3	0.41	0.52	0.46	25
4	0.43	0.57	0.49	21
5	0.57	0.72	0.63	18
6	0.72	0.50	0.59	26
7	0.58	0.55	0.56	20
8	0.75	0.75	0.75	16
9	0.80	0.64	0.71	25
10	0.41	0.45	0.43	20
11	0.48	0.67	0.56	24
12	0.44	0.34	0.39	32
13	0.22	0.21	0.21	24
14	0.52	0.65	0.58	23
15	0.61	0.50	0.55	22
16	0.80	0.73	0.76	22
17	0.29	0.31	0.30	26
18	0.26	0.23	0.24	22
19	0.26	0.26	0.26	23
20	0.24	0.30	0.27	20
21	0.52	0.38	0.44	34
22	0.38	0.15	0.21	20
accuracy				0.49
macro avg				0.49
weighted avg				0.48

Increase accuracy from 0.440 to 0.486 after tuning

Part II Advanced Model: Densely-connected Neural Networks

Procedure

BatchNorm -> Densely-connected NN -> ReLu -> Dropout -> BatchNorm -> Densely-connected NN -> ReLu -> Dropout -> Densely-connected NN -> ReLu -> Dropout -> Densely-connected NN -> ReLu -> Densely-connected NN -> Softmax -> Output

1. Provide directories for training/testing images.

```
In [3]: """  
        Path  
        """  
DATA_PATH = "../data/train_set"  
POINTS_FOLDER = os.path.join(DATA_PATH, "points")  
LABELS_FOLDER = DATA_PATH
```

2. Train/Test Split Feature Extraction

```

In [4]: def read_labels():
    labels_df = pd.read_csv(os.path.join(LABELS_FOLDER, 'label.csv'))
    labels_df = labels_df.loc[:, ['emotion_idx', 'emotion_cat', 'type']]
    return labels_df

def read_points():
    files = [file for file in os.listdir(POINTS_FOLDER) if file.endswith('.mat')]
    files.sort()

    face_points = np.zeros((len(files), 78, 2))
    for index, filename in enumerate(files):
        face_points_dict = loadmat(os.path.join(POINTS_FOLDER, filename))

        face_points[index] = face_points_dict.get('faceCoordinatesUnwarped',
        face_points_dict.get('faceCoordinates2'))
    return face_points

points = read_points()
labels = read_labels()

### train test split
X_points_train, X_points_test, y_train, y_test = train_test_split(points, labels, test_size=0.2, random_state=666)

### Feature Extraction time on training set:
feature_training_start = time.time()
X_train = np.zeros((X_points_train.shape[0], 3003))
for i in range(X_points_train.shape[0]):
    current = X_points_train[i]
    X_train[i,] = pdist(current)
feature_training_end = time.time()
y_train = y_train['emotion_idx']
y_train = to_categorical(y_train)[:,:1:]
print("Feature Extraction time on training set:", "%s seconds" % (feature_training_end - feature_training_start))

### Feature Extraction time on test set:
feature_test_start = time.time()
X_test = np.zeros((X_points_test.shape[0], 3003))
for i in range(X_points_test.shape[0]):
    current = X_points_test[i]
    X_test[i,] = pdist(current)
feature_test_end = time.time()
y_test = y_test['emotion_idx']
y_test = to_categorical(y_test)[:,:1:]
print("Feature Extraction time on test set:", "%s seconds" % (feature_test_end - feature_test_start))

```

Feature Extraction time on training set: 0.0984201431274414 seconds
 Feature Extraction time on test set: 0.026224851608276367 seconds

3. Train model


```
In [5]: input_shape = [3003]
input_layer = Input(input_shape)
x = BatchNormalization(momentum = 0.88)(input_layer)
x = Dense(22*10,activation='relu',kernel_initializer=initializers.glorot_normal(seed=4))(x)
x = Dropout(0.25)(x)
x = BatchNormalization()(x)
x = Dense(22*8,activation='relu',kernel_initializer=initializers.glorot_normal(seed=4))(x)
x = Dropout(0.25)(x)
x = Dense(22*4,activation='relu',kernel_initializer=initializers.glorot_normal(seed=4))(x)
x = Dropout(0.25)(x)
x = Dense(22*2,activation='relu',kernel_initializer=initializers.glorot_normal(seed=4))(x)
output_layer = Dense(22,activation='softmax',kernel_initializer=initializers.glorot_normal(seed=4))(x)
model2 = Model(input_layer,output_layer)
```

```
In [6]: start_time = time.time()
model2.compile(loss='categorical_crossentropy',optimizer = Adam(lr=0.001),metrics=['accuracy'])
model_history = model2.fit(X_train,y_train,epochs = 40,validation_data=[X_test,y_test])
print("training model takes %s seconds" % round((time.time() - start_time),3))
```

Train on 2000 samples, validate on 500 samples

Epoch 1/40

2000/2000 [=====] - 11s 6ms/step - loss: 3.0909 - accuracy: 0.0745 - val_loss: 2.8831 - val_accuracy: 0.1280

Epoch 2/40

2000/2000 [=====] - 3s 2ms/step - loss: 2.7894 - accuracy: 0.1450 - val_loss: 2.5699 - val_accuracy: 0.1820

Epoch 3/40

2000/2000 [=====] - 3s 2ms/step - loss: 2.4786 - accuracy: 0.2215 - val_loss: 2.1691 - val_accuracy: 0.2880

Epoch 4/40

2000/2000 [=====] - 4s 2ms/step - loss: 2.2981 - accuracy: 0.2600 - val_loss: 2.0427 - val_accuracy: 0.3080

Epoch 5/40

2000/2000 [=====] - 4s 2ms/step - loss: 2.1345 - accuracy: 0.2895 - val_loss: 1.8898 - val_accuracy: 0.3800

Epoch 6/40

2000/2000 [=====] - 4s 2ms/step - loss: 2.0055 - accuracy: 0.3040 - val_loss: 1.8210 - val_accuracy: 0.3620

Epoch 7/40

2000/2000 [=====] - 4s 2ms/step - loss: 1.9362 - accuracy: 0.3310 - val_loss: 1.7562 - val_accuracy: 0.3840

Epoch 8/40

2000/2000 [=====] - 4s 2ms/step - loss: 1.8543 - accuracy: 0.3705 - val_loss: 1.6571 - val_accuracy: 0.4320

Epoch 9/40

2000/2000 [=====] - 3s 2ms/step - loss: 1.7891 - accuracy: 0.3830 - val_loss: 1.6187 - val_accuracy: 0.4100

Epoch 10/40

2000/2000 [=====] - 4s 2ms/step - loss: 1.8106 - accuracy: 0.3845 - val_loss: 1.6062 - val_accuracy: 0.4460

Epoch 11/40

2000/2000 [=====] - 4s 2ms/step - loss: 1.6791 - accuracy: 0.4455 - val_loss: 1.5587 - val_accuracy: 0.4740

Epoch 12/40

2000/2000 [=====] - 4s 2ms/step - loss: 1.6583 - accuracy: 0.4335 - val_loss: 1.5786 - val_accuracy: 0.4380

Epoch 13/40

2000/2000 [=====] - 3s 2ms/step - loss: 1.6121 - accuracy: 0.4470 - val_loss: 1.5002 - val_accuracy: 0.4920

Epoch 14/40

2000/2000 [=====] - 4s 2ms/step - loss: 1.5910 - accuracy: 0.4515 - val_loss: 1.5508 - val_accuracy: 0.4800

Epoch 15/40

2000/2000 [=====] - 4s 2ms/step - loss: 1.5545 - accuracy: 0.4780 - val_loss: 1.5120 - val_accuracy: 0.5020

Epoch 16/40

2000/2000 [=====] - 3s 2ms/step - loss: 1.5371 - accuracy: 0.4740 - val_loss: 1.4611 - val_accuracy: 0.4800

Epoch 17/40

2000/2000 [=====] - 4s 2ms/step - loss: 1.5343 - accuracy: 0.4810 - val_loss: 1.4801 - val_accuracy: 0.5060

Epoch 18/40

2000/2000 [=====] - 4s 2ms/step - loss: 1.4862 - accuracy: 0.4930 - val_loss: 1.4538 - val_accuracy: 0.5200

Epoch 19/40

2000/2000 [=====] - 3s 2ms/step - loss: 1.4885 - acc

uracy: 0.4935 - val_loss: 1.4841 - val_accuracy: 0.5020
Epoch 20/40
2000/2000 [=====] - 3s 1ms/step - loss: 1.4582 - accuracy: 0.5155 - val_loss: 1.4644 - val_accuracy: 0.4960
Epoch 21/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.4012 - accuracy: 0.5200 - val_loss: 1.5052 - val_accuracy: 0.4720
Epoch 22/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.3938 - accuracy: 0.5155 - val_loss: 1.4515 - val_accuracy: 0.5200
Epoch 23/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.3690 - accuracy: 0.5350 - val_loss: 1.3926 - val_accuracy: 0.5120
Epoch 24/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.3682 - accuracy: 0.5395 - val_loss: 1.3942 - val_accuracy: 0.5260
Epoch 25/40
2000/2000 [=====] - 5s 2ms/step - loss: 1.3361 - accuracy: 0.5460 - val_loss: 1.4021 - val_accuracy: 0.5360
Epoch 26/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.3434 - accuracy: 0.5355 - val_loss: 1.4091 - val_accuracy: 0.5220
Epoch 27/40
2000/2000 [=====] - 3s 2ms/step - loss: 1.3059 - accuracy: 0.5390 - val_loss: 1.4268 - val_accuracy: 0.5260
Epoch 28/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.2742 - accuracy: 0.5600 - val_loss: 1.4122 - val_accuracy: 0.5260
Epoch 29/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.2681 - accuracy: 0.5700 - val_loss: 1.3886 - val_accuracy: 0.5200
Epoch 30/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.2435 - accuracy: 0.5820 - val_loss: 1.4358 - val_accuracy: 0.5220
Epoch 31/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.2800 - accuracy: 0.5715 - val_loss: 1.3830 - val_accuracy: 0.5280
Epoch 32/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.2505 - accuracy: 0.5670 - val_loss: 1.3526 - val_accuracy: 0.5280
Epoch 33/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.2514 - accuracy: 0.5840 - val_loss: 1.3533 - val_accuracy: 0.5260
Epoch 34/40
2000/2000 [=====] - 3s 1ms/step - loss: 1.1825 - accuracy: 0.5890 - val_loss: 1.3739 - val_accuracy: 0.5300
Epoch 35/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.2138 - accuracy: 0.5755 - val_loss: 1.4148 - val_accuracy: 0.5480
Epoch 36/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.1915 - accuracy: 0.5885 - val_loss: 1.3319 - val_accuracy: 0.5320
Epoch 37/40
2000/2000 [=====] - 3s 2ms/step - loss: 1.1862 - accuracy: 0.5905 - val_loss: 1.3199 - val_accuracy: 0.5440
Epoch 38/40
2000/2000 [=====] - 3s 2ms/step - loss: 1.1417 - acc

```

uracy: 0.6035 - val_loss: 1.3572 - val_accuracy: 0.5440
Epoch 39/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.1625 - acc
uracy: 0.5970 - val_loss: 1.4254 - val_accuracy: 0.5200
Epoch 40/40
2000/2000 [=====] - 4s 2ms/step - loss: 1.1474 - acc
uracy: 0.6155 - val_loss: 1.4307 - val_accuracy: 0.5400
training model takes 158.948 seconds

```

```


In [7]: fig, ax = plt.subplots(figsize=[12,8])
ax.plot(model_history.history['accuracy'],'r',linewidth=1.0, label = 'Training
Accuracy')
ax.legend()
plt.xlabel('Epochs ',fontsize=16)
plt.ylabel('Accuracy',fontsize=16)
plt.title('Accuracy Curves',fontsize=16)

```

```

Out[7]: Text(0.5, 1.0, 'Accuracy Curves')

```



4. Test accuracy

```

In [8]: preds = model2.evaluate(X_test, y_test)
print ("Loss = " + str(preds[0]))
print ("Test Accuracy = " + str(preds[1]))

```

```

500/500 [=====] - 0s 456us/step
Loss = 1.4307443332672118
Test Accuracy = 0.5400000214576721

```

5. Predicted label

```
In [9]: predict_DATA_PATH = "../data/test_set_predict"
POINTS_FOLDER = os.path.join(predict_DATA_PATH, "points")

predict_points = read_points()

predict_start = time.time()
X_predict = np.zeros((predict_points.shape[0], 3003))
for i in range(predict_points.shape[0]):
    current = predict_points[i]
    X_predict[i,] = pdist(current)
predict_end = time.time()
print("Feature Extraction time on test set:", "%s seconds"%(predict_end - predict_start))
```

Feature Extraction time on test set: 0.17246484756469727 seconds

```
In [27]: y_predict = []
for i in model2.predict(X_predict):
    y_predict.append(np.argmax(i) + 1)
```

```
In [12]: df_advance = pd.DataFrame(data = y_predict)
df_advance.to_csv('../output/baseline_prediction.csv')
```

```
In [25]: df_baseline = pd.DataFrame(data = baseline_tune.predict(X_predict))
df_baseline.to_csv('../output/baseline_prediction1.csv')
```